

Problem 1: Linear Regression from Scratch (30 points)

```
In [1]: # import the necessary packages
import numpy as np
from matplotlib import pyplot as plt
np.random.seed(100)
```

Let's generate some data points first, by the equation $y = x - 3$.


```
In [2]: x = np.random.randint(100, size=100)/30 - 2
X = x.reshape(-1, 1)

y = x + -3 + 0.3*np.random.randn(100)
```

Let's then visualize the data points we just created.

```
In [3]: plt.scatter(X, y)
```

```
Out[3]: <matplotlib.collections.PathCollection at 0x18a37e2a940>
```



1.1 Gradient of vanilla linear regression model (5 points)

In the lecture, we learn that the cost function of a linear regression model can be expressed as **Equation 1**:

$$J(\theta) = \frac{1}{2m} \sum_i^m \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right)^2$$

The gradient of it can be written as **Equation 2**:

$$\frac{\partial J(\theta)}{\partial \theta} = \frac{1}{m} \sum_i^m \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right) x^{(i)}$$

1.2 Gradient of vanilla regularized regression model (5 points)

After adding the L2 regularization term, the linear regression model can be expressed as

Equation 3:

$$J(\theta) = \frac{1}{2m} \sum_i^m \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right)^2 + \frac{\lambda}{2m} \sum_j^n (\theta_j)^2$$

The gradient of it can be written as **Equation 4:**

$$\frac{\partial J(\theta)}{\partial \theta} = \frac{1}{m} \sum_i^m \left(h_{\theta} \left(x^{(i)} \right) - y^{(i)} \right) x^{(i)} + \frac{\lambda}{m} (\theta)$$

1.3 Implement the cost function of a regularized regression model (5 points)

Please implement the cost function of a regularized regression model according to the above equations.

1.4 Implement the gradient of the cost function of a regularized regression model (5 points)

Please implement the gradient of the cost function of a regularized regression model according to the above equations.

```
In [16]: def regularized_linear_regression(X, y, alpha=0.01, lambda_value=1, epochs=30):
        """
        :param x: feature matrix
        :param y: target vector
        :param alpha: learning rate (default:0.01)
        :param lambda_value: lambda (default:1)
        :param epochs: maximum number of iterations of the
            linear regression algorithm for a single run (default=30)
        :return: weights, list of the cost function changing overtime
        """

        m = np.shape(X)[0] # total number of samples
        n = np.shape(X)[1] # total number of features

        X = np.concatenate((np.ones((m, 1)), X), axis=1)
        W = np.random.randn(n + 1, )

        # stores the updates on the cost function (loss function)
```

```

cost_history_list = []

# iterate until the maximum number of epochs
for current_iteration in np.arange(epochs): # begin the process

    # compute the dot product between our feature 'X' and weight 'W'
    y_estimated = X.dot(W)

    # calculate the difference between the actual and predicted value
    error = y_estimated - y

#####
##### Begin of Question 1.3 #####
#####

    ##### Please write down your code here:####

    # calculate the cost (MSE) (Equation 1)

    #note: elementwise
    cost_without_regularization = ((error **2).sum()) / (2*m)

    ##### Please write down your code here:####

    # regularization term

    reg_term = lambda_value / (2*m) * (W**2).sum()

    # calculate the cost (MSE) + regularization term (Equation 3)
    cost_with_regularization = cost_without_regularization + reg_term

#####
##### End of Question 1.3 #####
#####

#####
##### Begin of Question 1.4 #####
#####

    ##### Please write down your code here:####

    # calculate the gradient of the cost function with regularization term (Equation 4)
    error2 = np.array([error, error]).T
    g_without_reg = (np.multiply(error2, X)).sum(axis = 0) / m
    g_reg = lambda_value / m * W
    gradient = g_without_reg + g_reg

    # Now we have to update our weights
    W = W - alpha * gradient

#####
##### End of Question 1.4 #####
#####

```

```

# keep track the cost as it changes in each iteration
cost_history_list.append(cost_with_regularization)

# Let's print out the cost
print(f"Cost with regularization: {cost_with_regularization}")
print(f"Mean square error: {cost_without_regularization}")

return W, cost_history_list

```

Run the following code to train your model.

Hint: If you have the correct code written above, the cost should be 0.5181222986588751 when $\lambda = 10$.

```


In [17]: lambda_list = [0, 10, 100, 1000, 10000]
for lambda_ in lambda_list:
    # calls regression function with different values of lambda
    weight, _ = regularized_linear_regression(X, y, alpha=0.01,
                                              lambda_value=lambda_, epochs=1000)

    fitted_line = np.dot(X, weight[1]) + weight[0]
    plt.scatter(X, y, label='data points')
    plt.plot(X, fitted_line, color='r', label='Fitted line')
    plt.xlabel("X")
    plt.ylabel("y")
    plt.title(f"Regression (lambda : {lambda_})")
    plt.legend()
    plt.show()

```


Cost with regularization: 0.05165921361463465

Mean square error: 0.05165921361463465




Cost with regularization: 0.5181225987615888


Mean square error: 0.08984032593751458




Cost with regularization: 2.79317248874003
Mean square error: 1.2785107553109407



Cost with regularization: 5.591464362606628
Mean square error: 4.946888025066497



Cost with regularization: 6.242695626933973
Mean square error: 6.161442583355813



1.5 Analyze your results (10 points)

According to the above figures, what's the best choice of λ ?

Why the regressed line turns to be flat as we increase λ ?

Your answer:

1. The best choice of λ is 0 according to the above figures
2. As we increases λ , we will get more loss if the model has higher coefficient. Hence, the model will be trained to use smaller coefficient, which makes the regressed line flat.

Problem 2: Getting familiar with PyTorch (30 points)

```
In [9]: import torch
import numpy as np
import mltools as ml
from matplotlib import pyplot as plt
```

1. Load the " data/curve80.txt " data set, and split it into 75% / 25% training/test. We will use degree=5 for all the polynomial features

```
In [10]: # Your code:
data = np.genfromtxt("data/curve80.txt")
X = data[:,0]
X = np.atleast_2d(X).T # code expects shape (M,N) so make sure it's 2-dimensional
Y = data[:,1] # doesn't matter for Y
Xtr,Xte,Ytr,Yte = ml.splitData(X,Y,0.75) # split data set 75/25

degree = 5
XtrP = ml.transforms.fpoly(Xtr, degree=degree, bias=False)
XtrP,params = ml.transforms.rescale(XtrP)
```

Transform numpy arrays to tensor. Make sure the XtrP_tensor has the shape of (60, 5) while Ytr_tensor has the shape of (60, 1). (5 points)

```
In [11]: XtrP_tensor = torch.from_numpy(XtrP)
Ytr_tensor = torch.from_numpy(Ytr)

XtrP_tensor = XtrP_tensor.float()
Ytr_tensor = Ytr_tensor.float()

Ytr_tensor = torch.unsqueeze(Ytr_tensor, 1)
```

1. Initialize our linear regressor. (5 points)

```
In [12]: linear_regressor = torch.nn.Linear(in_features=5, out_features=1)
```

1. Set up the criterion and optimizer

```
In [13]: criterion = torch.nn.MSELoss()
optimizer = torch.optim.SGD(linear_regressor.parameters(), lr=0.1)
epochs = 100000
```

1. Training the regressor using gradient descent. 10 points


```
In [14]: loss_record = []
for _ in range(epochs):
    optimizer.zero_grad() # set gradient to zero
    pred_y = linear_regressor(XtrP_tensor)
    loss = criterion(pred_y, Ytr_tensor) # calculate loss function

    loss.backward() # backpropagate gradient
    loss_record.append(loss.item())
    optimizer.step() # update the parameters in the linear regressor
```

1. Plot the loss v.s. epochs. Show the plot here. 5 points.

```
In [15]: plt.plot(range(epochs), (loss_record))
```

```
Out[15]: [<matplotlib.lines.Line2D at 0x1e8cde8d820>]
```

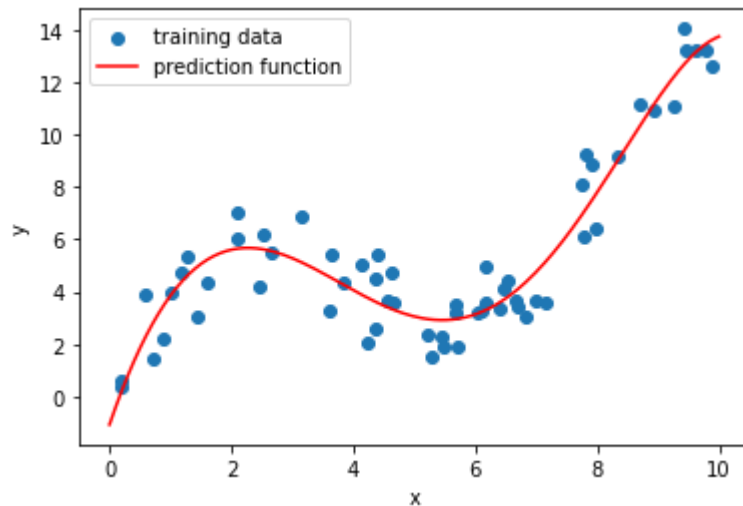


1. Visualize the trained linear regressor. 5 points.

```
In [16]: xs = np.linspace(0,10,200)
xs = xs[:,np.newaxis]
xsP, _ = ml.transforms.rescale(ml.transforms.fpoly(xs,degree=degree,bias=False), param
xsP_tensor = torch.from_numpy(xsP).float()
ys = linear_regressor(xsP_tensor)

plt.scatter(Xtr,Ytr,label="training data")
plt.plot(xs,ys.detach().numpy(),label="prediction function",color='red')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
```

Out[16]: <matplotlib.legend.Legend at 0x1e8ce830be0>



Statement of Collaboration

Sabina Yang - discussed about the usage of pytorch

In []: